

# greenspace\_portfolio

February 3, 2026

## 1 Set-Up

1.0.1 This code is a workflow expanding on previous work with Chicago census tracts and NDVI data. I am focussed on implementing income data and evaluating/refining the performance of the model.

```
[1]: ### Import libraries

import pathlib                      # file paths
import os                           # file paths
import warnings                     # warnings

import pandas as pd                 # dataframes
import geopandas as gpd            # geodataframes
import numpy as np                  # arrays
import xarray as xr                # multi-dimensional arrays
import rioxarray as rxr             # rasters + xarray
import rioxarray.merge as rxrm      # raster merging
import shapely                       # geometric operations
import time                          # time operations

import geoviews as gv               # geospatial visualizations
import holoviews as hv              # interactive visualizations
import hvplot.pandas                # pandas visualizations
import hvplot.xarray                 # xarray visualizations
from cartopy import crs as ccrs     # coordinate reference systems

import matplotlib                   # plotting
import matplotlib.pyplot as plt    # plotting
import scipy.stats as stats         # statistical analysis

import pystac_client                # STAC catalog access
from sodapy import Socrata          # Socrata Open Data API
import requests                     # JSON Census API requests

from tqdm.notebook import tqdm       # progress bars
from scipy.ndimage import convolve  # image processing
from scipy.ndimage import label     # image processing
```

```

from bokeh.models import NumeralTickFormatter # plot formatting

from sklearn.linear_model import LinearRegression          # linear regression
    ↵model
from sklearn.model_selection import train_test_split      # train-test split
from sklearn.model_selection import KFold                 # k-fold
    ↵cross-validation
from sklearn.metrics import mean_squared_error, r2_score  # model evaluation
    ↵metrics

```

[2]:

```

### Create reproducible file paths
dir = os.path.join(
    pathlib.Path.home(),

    'Data',
    'Earth Analytics',
    'greenspace-portfolio'

)

os.makedirs(dir, exist_ok=True)

```

[3]:

```

### Prevent GDAL from quitting due to momentary disruptions
os.environ["GDAL_HTTP_MAX_RETRY"] = "5"
os.environ["GDAL_HTTP_RETRY_DELAY"] = "1"

```

[4]:

```

### Set up the census tract path

tract_dir = os.path.join(dir, 'chicago-tract')
os.makedirs(tract_dir, exist_ok=True)
tract_path = os.path.join(tract_dir, 'chicago-tracts.shp')

### Download the census tracts (only once)
if not os.path.exists(tract_path):
    tract_url = ('https://data.cdc.gov/download/x7zy-2xmx/application%2Fzip')

    tract_gdf = gpd.read_file(tract_url)

    chi_tract_gdf = tract_gdf[tract_gdf.PlaceName == 'Chicago']

    chi_tract_gdf.to_file(tract_path, index = False)

### Load in census tract data

chi_tract_gdf = gpd.read_file(tract_path)
chi_tract_gdf

```

```

### Site plot -- Census tracts with satellite imagery in the background
(
    chi_tract_gdf

    .to_crs(ccrs.Mercator())

    .hvplot(
        line_color = 'blue', fill_color = None,
        crs = ccrs.Mercator(), tiles = 'EsriImagery',
        frame_width = 600,
        title = 'Chicago Census Tracts'
    )
)

```

[4]: :Overlay

```

.WMTS.I      :WMTS      [Longitude, Latitude]
.Polygons.I :Polygons   [Longitude, Latitude]

```

[5]: chi\_tract\_gdf

```

[5]:   place2010    tract2010  ST PlaceName      plctract10 PlcTrPop10 \
0     1714000  17031010100  17  Chicago  1714000-17031010100      4854
1     1714000  17031010201  17  Chicago  1714000-17031010201      6450
2     1714000  17031010202  17  Chicago  1714000-17031010202      2818
3     1714000  17031010300  17  Chicago  1714000-17031010300      6236
4     1714000  17031010400  17  Chicago  1714000-17031010400      5042
..
804    ...      ...      ...  ...
805    1714000  17031835700  17  Chicago  1714000-17031835700      0
806    1714000  17031980000  17  Chicago  1714000-17031980000      0
807    1714000  17043840000  17  Chicago  1714000-17043840000      0
808    1714000  17043840801  17  Chicago  1714000-17043840801      0

                           geometry
0     POLYGON ((-9758835.381 5164429.383, -9758837.3...
1     POLYGON ((-9760143.496 5163888.741, -9760143.4...
2     POLYGON ((-9759754.212 5163883.196, -9759726.6...
3     POLYGON ((-9758695.229 5163870.91, -9758695.78...
4     POLYGON ((-9757724.634 5160715.939, -9757742.2...
..
804    POLYGON ((-9754496.811 5132719.307, -9754491.8...
805    POLYGON ((-9788665.342 5161808.277, -9788667.6...
806    POLYGON ((-9766922.964 5128945.613, -9766938.5...
807    POLYGON ((-9787210.28 5154713.902, -9787240.67...
808    POLYGON ((-9787431.695 5154736.805, -9787441.0...

[809 rows x 7 columns]

```

```
[6]: ### Set up a path for the asthma data
cdc_path = os.path.join(dir, 'asthma.csv')

### Download asthma data (only once)
if not os.path.exists(cdc_path):
    cdc_url = (
        'https://data.cdc.gov/resource/cwsq-ngmh.csv'
        "?year=2023"
        "&stateabbr=IL"
        "&countyname=Cook"
        "&measureid=CASTHMA"
        "&$limit=1500"
    )

    cdc_df = (
        pd.read_csv(cdc_url)

        .rename(columns = {
            'data_value': 'asthma',
            'low_confidence_limit': 'asthma_ci_low',
            'high_confidence_limit': 'asthma_ci_high',
            'locationname': 'tract'})[[

            'year', 'tract', 'asthma',
            'asthma_ci_low', 'asthma_ci_high', 'data_value_unit',
            'totalpopulation', 'totalpop18plus',
        ]])
    )

    cdc_df.to_csv(cdc_path, index = False)

### Load in asthma data
cdc_df = pd.read_csv(cdc_path)

### Preview asthma data
cdc_df
```

	year	tract	asthma	asthma_ci_low	asthma_ci_high	\
0	2023	17031071700	9.1	8.1	10.1	
1	2023	17031120300	8.8	7.8	9.8	
2	2023	17031010501	9.9	8.9	11.0	
3	2023	17031031000	8.9	7.9	9.9	
4	2023	17031062400	9.2	8.2	10.3	
...	...	...	...	...	...	
1323	2023	17031840400	7.9	7.1	8.8	
1324	2023	17031837300	13.7	12.3	15.0	

```

1325 2023 17031840000      8.4          7.5          9.3
1326 2023 17031840700      8.3          7.3          9.2
1327 2023 17031838700     14.9         13.3         16.6

```

		data_value_unit	totalpopulation	totalpop18plus
0		%	1660	1325
1		%	6920	5212
2		%	4206	3762
3		%	3868	3439
4		%	1673	1389
...	...	...	...	...
1323		%	3369	2662
1324		%	2489	1835
1325		%	3001	2428
1326		%	3900	2885
1327		%	4132	2874

[1328 rows x 8 columns]

From what I can tell, it appears int64 and similar numerics will drop leading zeroes. Not an issue here, but something I wanted to try to avoid. I also found it much easier to convert the column types outside of the if statements downloading the data.

```

[7]: # Change tract identifier datatype for merging
cdc_df["tract"] = cdc_df["tract"].astype(str)
chi_tract_gdf["tract2010"] = chi_tract_gdf["tract2010"].astype(str)

# Merge census data with geometry
tract_cdc_gdf = (
    chi_tract_gdf.merge(
        cdc_df,
        left_on = 'tract2010',
        right_on = 'tract',
        how = 'inner'           # Only keep tracts with asthma data
    )
    [['tract2010', 'asthma', 'asthma_ci_low', 'asthma_ci_high', 'geometry']]
)
tract_cdc_gdf

```

```

[7]:      tract2010  asthma  asthma_ci_low  asthma_ci_high \
0   17031010100    11.3        10.1        12.5
1   17031010201    11.2        10.0        12.4
2   17031010202    10.2        9.1         11.3
3   17031010300    10.2        9.1         11.3
4   17031010400    10.6        9.5         11.7
...
783 17031770700     9.7        8.6         10.8

```

```

784 17031770800      9.5          8.4          10.6
785 17031805600      9.1          8.1          10.1
786 17031807900      8.5          7.5          9.4
787 17031808100      9.2          8.1          10.3

                           geometry
0   POLYGON ((-9758835.381 5164429.383, -9758837.3...
1   POLYGON ((-9760143.496 5163888.741, -9760143.4...
2   POLYGON ((-9759754.212 5163883.196, -9759726.6...
3   POLYGON ((-9758695.229 5163870.91, -9758695.78...
4   POLYGON ((-9757724.634 5160715.939, -9757742.2...
..
..                         ...
783  POLYGON ((-9780753.304 5157066.079, -9780752.0...
784  POLYGON ((-9783235.84 5154620.343, -9783211.23...
785  POLYGON ((-9776210.02 5161605.738, -9776213.47...
786  POLYGON ((-9768609.902 5160576.634, -9768654.5...
787  MULTIPOLYGON (((-9774480.671 5161127.722, -977...

```

[788 rows x 5 columns]

```
[8]: ### Plot asthma data as chloropleth
asthma_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        tract_cdc_gdf.to_crs(ccrs.Mercator()),
        vdims = ['asthma', 'tract2010'],
        crs = ccrs.Mercator()

        ).opts(color = 'asthma', colorbar = True, tools = ['hover'])
    ).opts(width = 600, height = 600, xaxis = None, yaxis = None)

asthma_plot
```

```
[8]: :Overlay
.WMTS.I      :WMTS      [Longitude, Latitude]
.Polygons.I :Polygons    [Longitude, Latitude]    (asthma, tract2010)
```

## 1.1 Access income data and add to the GDF

```
[9]: # Census API Key
income_API_key = '87f7d00c90e94cc038e9714efd302a96d6d49803'

# Set up a path for the income data
income_path = os.path.join(dir, 'income.csv')

# Download income data (only once)
```

```

if not os.path.exists(income_path):
    income_url = "https://api.census.gov/data/2023/acs/acs5" # Income only
    ↪tracked on ACS 5-Year Census data
    parameters = {
        "get": "B19013_001E,B19301_001E", # 'get' median household income and
        ↪per-capita income respectively
        "for": "tract:*",                 # one row per census tract
        "in": "state:17 county:031",      # 'in' Cook County, Illinois
        "key": income_API_key
    }

    # Make the API request
    response = requests.get(income_url, params = parameters)
    income_data = response.json()

    # Convert to DF
    income_df = pd.DataFrame(
        income_data[1:],
        columns = income_data[0]
    )

    # # Progress check
    # print(income_df.head())

    income_df = income_df.rename(columns={
        "B19013_001E": "median_household_income",
        "B19301_001E": "per_capita_income"
    })

    # Convert income columns to numeric
    income_df["median_household_income"] = pd.to_numeric(
        income_df["median_household_income"], errors="coerce")

    income_df["per_capita_income"] = pd.to_numeric(
        income_df["per_capita_income"], errors="coerce")

    # Create tractID column for merging
    income_df["tractID"] = income_df["state"] + income_df["county"] + ↪
    ↪income_df["tract"]

    # # Progress check
    # print(income_df.head())

# Save to CSV
else:
    income_df = pd.read_csv(income_path)

```

```
income_df
```

```
[9]: median_household_income per_capita_income state county tract \
0 69460 45353 17 31 10100
1 49639 31978 17 31 10201
2 55119 33488 17 31 10202
3 65871 42487 17 31 10300
4 49017 33185 17 31 10400
...
1327 56848 29766 17 31 844600
1328 69583 27809 17 31 844700
1329 -666666666 17 31 980000
1330 -666666666 17 31 980100
1331 -666666666 17 31 990000

tractID
0 17031010100
1 17031010201
2 17031010202
3 17031010300
4 17031010400
...
1327 17031844600
1328 17031844700
1329 17031980000
1330 17031980100
1331 17031990000
```

[1332 rows x 6 columns]

```
[10]: # Change tract identifier datatype for merging
income_df["tractID"] = income_df["tractID"].astype(str)

# Combine income data with existing tract + asthma data
inc_tract_df = (
    income_df.merge(
        tract_cdc_gdf,
        left_on = "tractID",
        right_on = "tract2010",
        how = "inner" # Only keep tracts with full data
    )
    [['tractID', 'asthma', 'asthma_ci_low', 'asthma_ci_high', 'median_household_income', 'per_capita_income', 'geometry', ]]
```

```
inc_tract_df
```

[10]:

	tractID	asthma	asthma_ci_low	asthma_ci_high	\
0	17031010100	11.3	10.1	12.5	
1	17031010201	11.2	10.0	12.4	
2	17031010202	10.2	9.1	11.3	
3	17031010300	10.2	9.1	11.3	
4	17031010400	10.6	9.5	11.7	
..	..	..	..	..	
783	17031843500	9.5	8.3	10.7	
784	17031843600	12.5	11.2	14.0	
785	17031843700	8.9	7.9	9.9	
786	17031843800	10.9	9.7	12.0	
787	17031843900	13.1	11.7	14.7	
	median_household_income	per_capita_income	\		
0	69460	45353			
1	49639	31978			
2	55119	33488			
3	65871	42487			
4	49017	33185			
..	..	..			
783	55833	3546			
784	71359	38837			
785	134000	101400			
786	57885	27767			
787	47760	34298			
	geometry				
0	POLYGON ((-9758835.381 5164429.383, -9758837.3...				
1	POLYGON ((-9760143.496 5163888.741, -9760143.4...				
2	POLYGON ((-9759754.212 5163883.196, -9759726.6...				
3	POLYGON ((-9758695.229 5163870.91, -9758695.78...				
4	POLYGON ((-9757724.634 5160715.939, -9757742.2...				
..	..	..			
783	POLYGON ((-9761322.813 5136241.389, -9761337.5...				
784	POLYGON ((-9752318.844 5132492.145, -9752318.2...				
785	POLYGON ((-9761412.872 5153004.52, -9761411.75...				
786	POLYGON ((-9755524.513 5130243.96, -9755529.18...				
787	POLYGON ((-9748410.585 5128192.165, -9748480.7...				

[788 rows x 7 columns]

[11]: # Convert to GeoDataFrame

```
inc_tract_gdf = gpd.GeoDataFrame(
    inc_tract_df,
    geometry='geometry',
```

```

        crs=chi_tract_gdf.crs # use the original CRS from your tract GDF
    )

[12]: ### Plot income data as chloropleth
pc_inc_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        inc_tract_gdf.to_crs(ccrs.Mercator()),
        vdims = ['per_capita_income', 'tractID'],
        crs = ccrs.Mercator()

        ).opts(color = 'per_capita_income', colorbar = True, tools = ['hover'],
               colorbar_opts = {'formatter': NumeralTickFormatter(format="$0,0")}) ↴
        # Format colorbar as currency
    ).opts(width = 600, height = 600, xaxis = None, yaxis = None)

pc_inc_plot

[12]: :Overlay
.WMTS.I      :WMTS [Longitude,Latitude]
.Polygons.I :Polygons [Longitude,Latitude] (per_capita_income,tractID)

[13]: ### Plot median income data as chloropleth
med_inc_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        inc_tract_gdf.to_crs(ccrs.Mercator()),
        vdims = ['median_household_income', 'tractID'],
        crs = ccrs.Mercator()

        ).opts(color = 'median_household_income', colorbar = True, tools = [
            'hover'],
               colorbar_opts = {'formatter': NumeralTickFormatter(format="$0,0")}) ↴
        # Format colorbar as currency
    ).opts(width = 600, height = 600, xaxis = None, yaxis = None)

med_inc_plot

[13]: :Overlay
.WMTS.I      :WMTS [Longitude,Latitude]
.Polygons.I :Polygons [Longitude,Latitude]
(median_household_income,tractID)

[14]: # Replace extreme negative outliers with NaN

```

```

inc_tract_gdf['median_household_income'] =_
    ↪inc_tract_gdf['median_household_income'].replace({-666666666: np.nan})
inc_tract_gdf['per_capita_income'] = inc_tract_gdf['per_capita_income']._
    ↪replace({-666666666: np.nan})

```

[15]: *### Plot median income data as chloropleth*

```

med_inc_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        inc_tract_gdf.to_crs(ccrs.Mercator()),
        vdims = ['median_household_income', 'tractID'],
        crs = ccrs.Mercator()

        ).opts(color = 'median_household_income', colorbar = True, tools =_
            ↪['hover'],
            colorbar_opts = {'formatter': NumeralTickFormatter(format="$0,0")})_
            ↪    # Format colorbar as currency
        ).opts(width = 600, height = 600, xaxis = None, yaxis = None)

med_inc_plot

```

[15]: *:Overlay*

```

.WMTS.I      :WMTS      [Longitude, Latitude]
.Polygons.I :Polygons   [Longitude, Latitude]
(median_household_income, tractID)

```

## 1.2 NDVI Data

[16]: *### Connect to the planetary computer catalog*

```

e84_catalog = pystac_client.Client.open(
    "https://planetarycomputer.microsoft.com/api/stac/v1"
)

```

[ ]: *# ONLY RUN ONCE TO DOWNLOAD DATA*

```

### Convert geometry to lat/lon for STAC
tract_latlong_gdf = inc_tract_gdf.to_crs(epsg = 4326)

### Define a path to save NDVI stats
ndvi_stats_path = os.path.join(dir, 'chicago-ndvi-stats.csv')

### Check for existing data - do not access duplicate tracts
downloaded_tracts = []

if os.path.exists(ndvi_stats_path):
    ndvi_stats_df = pd.read_csv(ndvi_stats_path)

```

```

downloaded_tracts = ndvi_stats_df.tract.values
else:
    print("No census tracts downloaded yet.")

### Loop through each census tract

scene_dfs = []

for i, tract_values in tqdm(tract_latlong_gdf.iterrows()):

    tract = tract_values.tractID

    ### Check if statistics are already downloaded for this tract
    if not (tract in downloaded_tracts):

        ### Repeat up to 5 times in case of a momentary disruption
        i = 0
        retry_limit = 5
        while i < retry_limit:

            ### Try accessing the STAC
            try:

                ### Search for tiles
                naip_search = e84_catalog.search(
                    collections = ['naip'],

                    intersects = shapely.to_geojson(tract_values.geometry),
                    datetime="2021"
                )

                ### Build dataframe with tracts and tile urls
                scene_dfs.append(pd.DataFrame(dict(
                    tract = tract,
                    date = [pd.to_datetime(scene.datetime).date()]

                    for scene in naip_search.items()]
                ,)

                rgbir_href = [scene.assets['image'].href for scene in
                ↪naip_search.items()]
            )))

            break

        ### Try again in case of an APIError

```

```

        except pystac_client.exceptions.APIError:
            print(
                f"Could not connect to STAC for tract {tract}. "
                f"Retrying tract {tract}..."
            )
            time.sleep(2)
            i+=1
            continue

    ### Concatenate the url dataframes
    if scene_dfs:
        scene_df = pd.concat(scene_dfs).reset_index(drop = True)
    else:
        scene_df = None

    ### Preview the url dataframe
    scene_df

```

[ ]: # Save to csv, will only work when downloading data for the first time

```

scene_df.to_csv(ndvi_stats_path, index = False)

```

[18]: # make csv for streaming ndvi data, has to be run each time to set the path

```

ndvi_stats_path_veg = os.path.join(dir, 'chicago-ndvi-veg-test.csv')

```

[ ]: # ONLY RUN ONCE TO DOWNLOAD DATA

```

downloaded_tracts = []

if os.path.exists(ndvi_stats_path):
    ndvi_stats_df = pd.read_csv(ndvi_stats_path)
    downloaded_tracts = ndvi_stats_df.tract.values # already processed tracts

# Filter scene_df to only tracts that haven't been processed
if scene_df is not None:
    scene_df_to_process = scene_df[~scene_df['tract'].isin(downloaded_tracts)]
else:
    scene_df_to_process = None

# Full loop for all tracts
if scene_df_to_process is not None and len(scene_df_to_process) > 0:

    ### Loop through the census tracts with URLs
    for tract, tract_date_gdf in tqdm(scene_df_to_process.groupby('tract')):

```

```

#### Open all images for tract
tile_das = []

# loop through each scene for this tract
for _, href_s in tract_date_gdf.iterrows():

    #### Open vsi connection to data
    tile_da = rxr.open_rasterio(

        # location of the raster image
        href_s.rgbir_href,

        # deal with no data, remove extra dimensions
        masked = True).squeeze()

    #### Clip data
    boundary = (
        inc_tract_gdf

        # deal with integer issues
        .assign(tractID = lambda df: df["tractID"].astype(str))

        # use tract ID as index
        .set_index('tractID')

        # select the geometry for this tract
        .loc[[str(tract)]])

    #reproject to match the tile CRS
    .to_crs(tile_da.rio.crs)

    #extract the geometry
    .geometry)

    #### Crop to bounding box
    crop_da = tile_da.rio.clip_box(
        *boundary.envelope.total_bounds,
        auto_expand = True)

    #### Clip data to the boundary of the census tract
    clip_da = crop_da.rio.clip(
        boundary,
        all_touched = True)

    #### Compute NDVI
    ndvi_da = (
        (clip_da.sel(band = 4) - clip_da.sel(band = 1)) /

```

```

        (clip_da.sel(band = 4) + clip_da.sel(band = 1))
    )

    ### Accumulate result
tile_das.append(ndvi_da)

### Merge data
scene_da = rxrm.merge_arrays(tile_das)

### Mask vegetation
veg_mask = (scene_da > 0.3)

### Calculate statistics and save data to file

# count all valid pixels
total_pixels = scene_da.notnull().sum()

# count all vegetated pixels
veg_pixels = veg_mask.sum()

### identify vegetation patches
labeled_patches, num_patches = label(veg_mask)

# count patch pixels
patch_sizes = np.bincount(labeled_patches.ravel())[1:]

### Calculate mean patch size
mean_patch_size = patch_sizes.mean()

### Calculate edge density

### Make kernel to calculate edge density
kernel = np.array([
    [1, 1, 1],
    [1, -8, 1],
    [1, 1, 1]])

### detect boundaries
edges = convolve(veg_mask, kernel, mode='constant')

# calculate edge density
edge_density = np.sum(edges != 0) / veg_mask.size

### Add a row to the statistics file for this tract
# create a dataframe for this tract
pd.DataFrame(dict(

```

```

#store tract ID
tract = [tract],

#store total pixels
total_pixels = [int(total_pixels)],

#store fraction of vegetated pixels
fract_veg = [float(veg_pixels/total_pixels)],

#store mean patch size
mean_patch_size = [mean_patch_size],

#store edge density
edge_density = [edge_density]

# write out as csv and save
)).to_csv(
    ndvi_stats_path_veg,
    #append mode
    mode = 'a',
    # get rid of row numbers
    index = False,
    # write header only if file does not already exist
    header = (not os.path.exists(ndvi_stats_path_veg))
)

```

[19]:

```
# check the file is loading correctly
# Save to CSV
ndvi_stats_df = pd.read_csv(ndvi_stats_path_veg)

ndvi_stats_df
```

[19]:

	tract	total_pixels	fract_veg	mean_patch_size	edge_density
0	17031010100	1059033	0.178657	55.225919	0.118612
1	17031010201	1531554	0.213859	57.543394	0.161668
2	17031010202	978546	0.186055	63.260250	0.123673
3	17031010300	1308392	0.191675	57.113642	0.126384
4	17031010400	1516964	0.198563	52.983817	0.079474
..	...	...	...	...	...
783	17031843500	5647650	0.075254	9.732779	0.104686
784	17031843600	1142916	0.054393	9.177296	0.101217
785	17031843700	6025768	0.027644	7.477533	0.047642
786	17031843800	3639014	0.093920	24.169295	0.105052
787	17031843900	4521964	0.199113	23.985215	0.124282

```
[788 rows x 5 columns]
```

```
[20]: # Change tract identifier datatype for merging
ndvi_stats_df["tract"] = ndvi_stats_df["tract"].astype(str)

# Add NDVI statistics to existing dataframe
ndvi_inc_gdf = (
    inc_tract_gdf.merge(
        ndvi_stats_df,
        left_on = 'tractID',
        right_on = 'tract',
        how = 'inner'           # Only keep tracts with NDVI data
    )
)

ndvi_inc_gdf.head()
```

```
[20]:      tractID  asthma  asthma_ci_low  asthma_ci_high  \
0  17031010100    11.3          10.1          12.5
1  17031010201    11.2          10.0          12.4
2  17031010202    10.2          9.1          11.3
3  17031010300    10.2          9.1          11.3
4  17031010400    10.6          9.5          11.7

      median_household_income  per_capita_income  \
0                  69460.0            45353
1                  49639.0            31978
2                  55119.0            33488
3                  65871.0            42487
4                  49017.0            33185

                                         geometry      tract  \
0  POLYGON ((-9758835.381 5164429.383, -9758837.3...  17031010100
1  POLYGON ((-9760143.496 5163888.741, -9760143.4...  17031010201
2  POLYGON ((-9759754.212 5163883.196, -9759726.6...  17031010202
3  POLYGON ((-9758695.229 5163870.91, -9758695.78...  17031010300
4  POLYGON ((-9757724.634 5160715.939, -9757742.2...  17031010400

      total_pixels  fract_veg  mean_patch_size  edge_density
0         1059033    0.178657      55.225919     0.118612
1         1531554    0.213859      57.543394     0.161668
2          978546    0.186055      63.260250     0.123673
3         1308392    0.191675      57.113642     0.126384
4         1516964    0.198563      52.983817     0.079474
```

```
[21]: fract_veg_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        ndvi_inc_gdf.to_crs(ccrs.Mercator()),
        vdims = ['fract_veg', 'tractID'],
        crs = ccrs.Mercator()

    ).opts(color = 'fract_veg', colorbar = True, tools = ['hover'])
).opts(width = 600, height = 600, xaxis = None, yaxis = None)

fract_veg_plot
```

```
[21]: :Overlay
.WMTS.I      :WMTS      [Longitude, Latitude]
.Polygons.I :Polygons    [Longitude, Latitude]    (fract_veg, tractID)
```

```
[22]: mean_veg_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(
        ndvi_inc_gdf.to_crs(ccrs.Mercator()),
        vdims = ['mean_patch_size', 'tractID'],
        crs = ccrs.Mercator()

    ).opts(color = 'mean_patch_size', colorbar = True, tools = ['hover'])
).opts(width = 600, height = 600, xaxis = None, yaxis = None)

mean_veg_plot
```

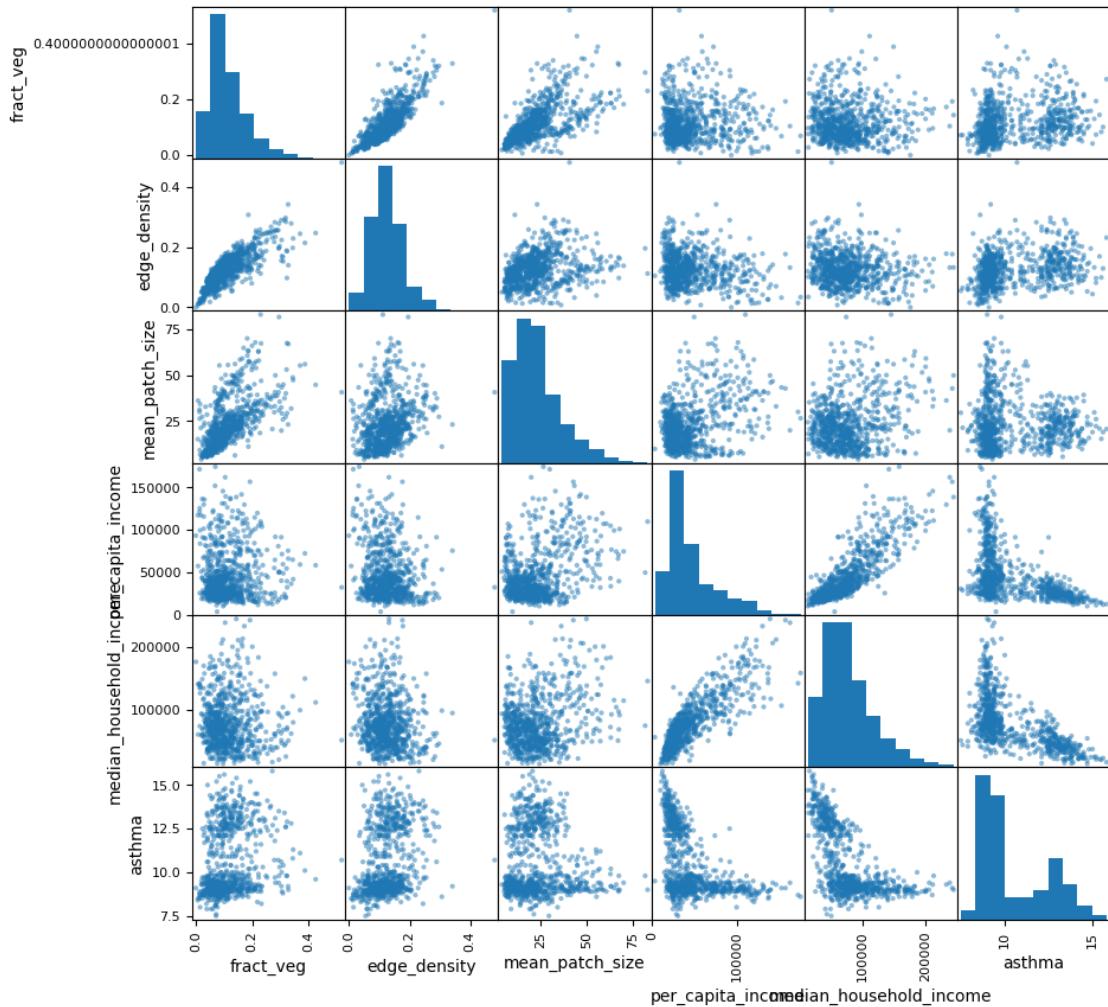
```
[22]: :Overlay
.WMTS.I      :WMTS      [Longitude, Latitude]
.Polygons.I :Polygons    [Longitude, Latitude]    (mean_patch_size, tractID)
```

[23]: *### Visualize distribution of data*

```
# graph for variables
vars = ['fract_veg', 'edge_density', 'mean_patch_size',
        'per_capita_income', 'median_household_income', 'asthma']

# scatter plot matrix
pd.plotting.scatter_matrix(
    ndvi_inc_gdf[vars],
    figsize = (10, 10)
)
```

```
[23]: array([[<Axes: xlabel='fract_veg', ylabel='fract_veg'>,
   <Axes: xlabel='edge_density', ylabel='fract_veg'>,
   <Axes: xlabel='mean_patch_size', ylabel='fract_veg'>,
   <Axes: xlabel='per_capita_income', ylabel='fract_veg'>,
   <Axes: xlabel='median_household_income', ylabel='fract_veg'>,
   <Axes: xlabel='asthma', ylabel='fract_veg'>],
  [<Axes: xlabel='fract_veg', ylabel='edge_density'>,
   <Axes: xlabel='edge_density', ylabel='edge_density'>,
   <Axes: xlabel='mean_patch_size', ylabel='edge_density'>,
   <Axes: xlabel='per_capita_income', ylabel='edge_density'>,
   <Axes: xlabel='median_household_income', ylabel='edge_density'>,
   <Axes: xlabel='asthma', ylabel='edge_density'>],
  [<Axes: xlabel='fract_veg', ylabel='mean_patch_size'>,
   <Axes: xlabel='edge_density', ylabel='mean_patch_size'>,
   <Axes: xlabel='mean_patch_size', ylabel='mean_patch_size'>,
   <Axes: xlabel='per_capita_income', ylabel='mean_patch_size'>,
   <Axes: xlabel='median_household_income', ylabel='mean_patch_size'>,
   <Axes: xlabel='asthma', ylabel='mean_patch_size'>],
  [<Axes: xlabel='fract_veg', ylabel='per_capita_income'>,
   <Axes: xlabel='edge_density', ylabel='per_capita_income'>,
   <Axes: xlabel='mean_patch_size', ylabel='per_capita_income'>,
   <Axes: xlabel='per_capita_income', ylabel='per_capita_income'>,
   <Axes: xlabel='median_household_income', ylabel='per_capita_income'>,
   <Axes: xlabel='asthma', ylabel='per_capita_income'>],
  [<Axes: xlabel='fract_veg', ylabel='median_household_income'>,
   <Axes: xlabel='edge_density', ylabel='median_household_income'>,
   <Axes: xlabel='mean_patch_size', ylabel='median_household_income'>,
   <Axes: xlabel='per_capita_income', ylabel='median_household_income'>,
   <Axes: xlabel='median_household_income',
   ylabel='median_household_income'>,
   <Axes: xlabel='asthma', ylabel='median_household_income'>],
  [<Axes: xlabel='fract_veg', ylabel='asthma'>,
   <Axes: xlabel='edge_density', ylabel='asthma'>,
   <Axes: xlabel='mean_patch_size', ylabel='asthma'>,
   <Axes: xlabel='per_capita_income', ylabel='asthma'>,
   <Axes: xlabel='median_household_income', ylabel='asthma'>,
   <Axes: xlabel='asthma', ylabel='asthma'>]], dtype=object)
```



```
[24]: # drop missing data

model_gdf = (
    ndvi_inc_gdf

    #make a copy to avoid modifying original data
    .copy()

    # subset to columns
    [['tractID', 'fract_veg', 'edge_density', 'mean_patch_size',
       'per_capita_income', 'median_household_income', 'asthma', 'geometry']]

    # drop NA
    .dropna()
)
```

```
model_gdf
```

```
[24]:      tractID  fract_veg  edge_density  mean_patch_size  per_capita_income  \
0    17031010100    0.178657    0.118612      55.225919        45353
1    17031010201    0.213859    0.161668      57.543394        31978
2    17031010202    0.186055    0.123673      63.260250        33488
3    17031010300    0.191675    0.126384      57.113642        42487
4    17031010400    0.198563    0.079474      52.983817        33185
..          ...
783   17031843500    0.075254    0.104686      9.732779        3546
784   17031843600    0.054393    0.101217      9.177296        38837
785   17031843700    0.027644    0.047642      7.477533       101400
786   17031843800    0.093920    0.105052      24.169295        27767
787   17031843900    0.199113    0.124282      23.985215        34298

      median_household_income  asthma  \
0                69460.0    11.3
1                49639.0    11.2
2                55119.0    10.2
3                65871.0    10.2
4                49017.0    10.6
..          ...
783              55833.0     9.5
784              71359.0    12.5
785            134000.0     8.9
786              57885.0    10.9
787              47760.0    13.1

      geometry
0    POLYGON ((-9758835.381 5164429.383, -9758837.3...
1    POLYGON ((-9760143.496 5163888.741, -9760143.4...
2    POLYGON ((-9759754.212 5163883.196, -9759726.6...
3    POLYGON ((-9758695.229 5163870.91, -9758695.78...
4    POLYGON ((-9757724.634 5160715.939, -9757742.2...
..          ...
783  POLYGON ((-9761322.813 5136241.389, -9761337.5...
784  POLYGON ((-9752318.844 5132492.145, -9752318.2...
785  POLYGON ((-9761412.872 5153004.52, -9761411.75...
786  POLYGON ((-9755524.513 5130243.96, -9755529.18...
787  POLYGON ((-9748410.585 5128192.165, -9748480.7...

[777 rows x 8 columns]
```

```
[25]: # Calculate logs
for col in ['fract_veg', 'edge_density', 'mean_patch_size',
            'per_capita_income', 'median_household_income', 'asthma']:
```

```

model_gdf[f'log_{col}'] = np.log(model_gdf[col])

model_gdf

[25]:      tractID  fract_veg  edge_density  mean_patch_size  per_capita_income \
0    17031010100  0.178657  0.118612  55.225919  45353
1    17031010201  0.213859  0.161668  57.543394  31978
2    17031010202  0.186055  0.123673  63.260250  33488
3    17031010300  0.191675  0.126384  57.113642  42487
4    17031010400  0.198563  0.079474  52.983817  33185
..
783   ...        ...
784   ...        ...
785   ...        ...
786   ...        ...
787   ...        ...

      median_household_income  asthma  \
0            69460.0       11.3
1            49639.0       11.2
2            55119.0       10.2
3            65871.0       10.2
4            49017.0       10.6
..
783   ...
784   ...
785   ...
786   ...
787   ...

      geometry  log_fract_veg  \
0  POLYGON ((-9758835.381 5164429.383, -9758837.3...  -1.722286
1  POLYGON ((-9760143.496 5163888.741, -9760143.4...  -1.542437
2  POLYGON ((-9759754.212 5163883.196, -9759726.6...  -1.681715
3  POLYGON ((-9758695.229 5163870.91, -9758695.78...  -1.651954
4  POLYGON ((-9757724.634 5160715.939, -9757742.2...  -1.616649
..
783   ...
784   ...
785   ...
786   ...
787   ...

      log_edge_density  log_mean_patch_size  log_per_capita_income  \
0            -2.131894          4.011432          10.722232
1            -1.822208          4.052539          10.372803
2            -2.090110          4.147257          10.418942

```

3	-2.068434	4.045043	10.656953
4	-2.532326	3.969987	10.409853
..	...	...	...
783	-2.256786	2.275499	8.173575
784	-2.290487	2.216733	10.567129
785	-3.044045	2.011903	11.526828
786	-2.253298	3.185083	10.231604
787	-2.085204	3.177438	10.442842
	log_median_household_income	log_asthma	
0	11.148506	2.424803	
1	10.812532	2.415914	
2	10.917250	2.322388	
3	11.095454	2.322388	
4	10.799922	2.360854	
..	...	...	
783	10.930120	2.251292	
784	11.175479	2.525729	
785	11.805595	2.186051	
786	10.966214	2.388763	
787	10.773944	2.572612	

[777 rows x 14 columns]

[26]: # q-q plots

```
# set variables for q-q plots
var_qq = [
    'fract_veg', 'log_fract_veg', 'edge_density', 'log_edge_density',
    'mean_patch_size', 'log_mean_patch_size',
    'per_capita_income', 'log_per_capita_income', 'median_household_income', 'log_median_household_income'
]

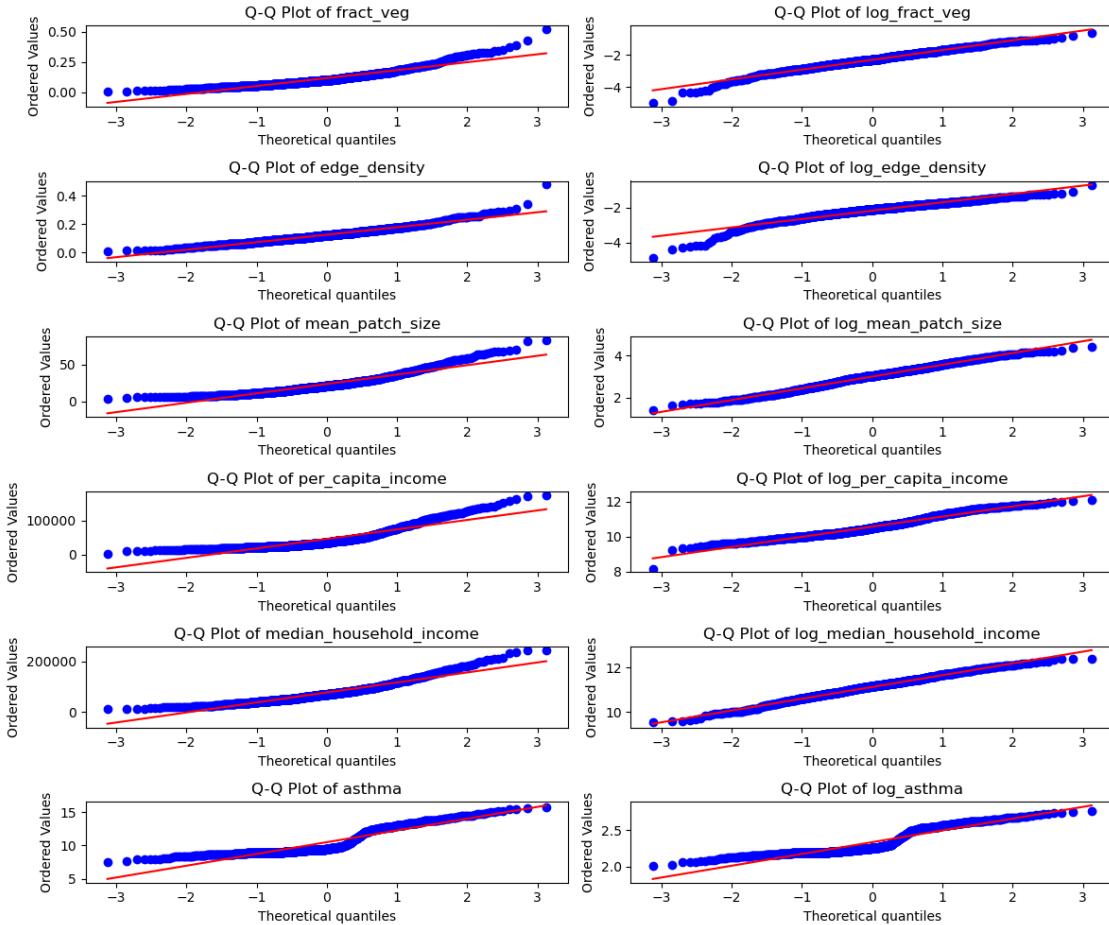
# make plot for each variable
plt.figure(figsize = (12,10))
for i, var in enumerate(var_qq, 1):

    # 2x2 facet
    plt.subplot(6, 2, i)

    # norm distribution q-q plot
    stats.probplot(model_gdf[var], dist="norm", plot=plt)

    # add title
    plt.title(f'Q-Q Plot of {var}')
```

```
plt.tight_layout()
plt.show()
```



From what I can tell from the histograms and q-q plots, fract\_veg and edge\_density as well as median\_household\_income and per\_capita\_income are correlated most closely to each other so I will pick one (or the log version) from each pair. Log\_fract\_veg seems to perform best of the top four options from non-exhaustive prior testing, and log\_median\_household\_income looks most normally distributed of that grouping. I will now build a model with log\_fract\_veg, log\_median\_household\_income, and log\_mean\_patch\_size as x variables and asthma as our y variable.

### 1.3 Run Regression

```
[29]: # Set predictor and response variables for modeling
X = model_gdf[['log_fract_veg', 'log_mean_patch_size', 'log_median_household_income']]
Y = model_gdf[['log_asthma']]
```

```

# Split data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.33, random_state=42
)

# Fit linear regression model
reg = LinearRegression()
reg.fit(X_train, Y_train)

# Make predictions on the test set
Y_test['pred_asthma'] = np.exp(reg.predict(X_test))

# use the trained model to predict asthma for all tracts
model_gdf['pred_asthma'] = np.exp(reg.predict(X))

Y_test['measured_asthma'] = np.exp(Y_test['log_asthma'])

Y_test

```

[29]:

	log_asthma	pred_asthma	measured_asthma
574	2.332144	11.898119	10.3
535	2.091864	10.189445	8.1
616	2.602690	11.847563	13.5
109	2.240710	10.113142	9.4
592	2.493205	12.858251	12.1
..	..	..	..
18	2.151762	9.718228	8.6
576	2.219203	11.059941	9.2
141	2.186051	7.897718	8.9
337	2.617396	11.558575	13.7
260	2.174752	8.782103	8.8

[257 rows x 3 columns]

[30]:

```

# vmin = Y_test[cols].min().min()
# vmax = Y_test[cols].max().max()

p_est = (Y_test
    .hvplot.scatter(x='measured_asthma', y='pred_asthma',
                    xlabel = 'Measured Asthma Prevalence',
                    ylabel = 'Predicted Asthma Prevalence',
                    title = 'Regression Model Distribution')
    .opts(aspect='equal',
          # xlim=(xylim), ylim=(xylim),
          width=600, height=600)
) * hv.Slope(slope=1, y_intercept=0).opts(color='black')

```

```
p_est
```

```
[30]: :Overlay
    .Scatter.I      :Scatter      [measured_asthma]      (pred_asthma)
    .Annotation.I :Slope       [x,y]
```

```
[31]: model_gdf = gpd.GeoDataFrame(
    model_gdf,
    geometry='geometry',
    crs=ndvi_inc_gdf.crs  # use the original CRS
)
```

```
[32]: model_gdf
```

```
[32]:      tractID  fract_veg  edge_density  mean_patch_size  per_capita_income \
0      17031010100  0.178657  0.118612  55.225919  45353
1      17031010201  0.213859  0.161668  57.543394  31978
2      17031010202  0.186055  0.123673  63.260250  33488
3      17031010300  0.191675  0.126384  57.113642  42487
4      17031010400  0.198563  0.079474  52.983817  33185
...
783    17031843500  0.075254  0.104686  9.732779  3546
784    17031843600  0.054393  0.101217  9.177296  38837
785    17031843700  0.027644  0.047642  7.477533  101400
786    17031843800  0.093920  0.105052  24.169295  27767
787    17031843900  0.199113  0.124282  23.985215  34298

      median_household_income  asthma  \
0            69460.0        11.3
1            49639.0        11.2
2            55119.0        10.2
3            65871.0        10.2
4            49017.0        10.6
...
783          55833.0        9.5
784          71359.0       12.5
785         134000.0        8.9
786          57885.0       10.9
787          47760.0       13.1

      geometry  log_fract_veg \
0  POLYGON ((-9758835.381 5164429.383, -9758837.3...  -1.722286
1  POLYGON ((-9760143.496 5163888.741, -9760143.4...  -1.542437
2  POLYGON ((-9759754.212 5163883.196, -9759726.6...  -1.681715
3  POLYGON ((-9758695.229 5163870.91, -9758695.78...  -1.651954
4  POLYGON ((-9757724.634 5160715.939, -9757742.2...  -1.616649
...
...
```

```

783 POLYGON ((-9761322.813 5136241.389, -9761337.5...      -2.586880
784 POLYGON ((-9752318.844 5132492.145, -9752318.2...      -2.911514
785 POLYGON ((-9761412.872 5153004.52, -9761411.75...      -3.588343
786 POLYGON ((-9755524.513 5130243.96, -9755529.18...      -2.365307
787 POLYGON ((-9748410.585 5128192.165, -9748480.7...      -1.613884

    log_edge_density  log_mean_patch_size  log_per_capita_income \
0            -2.131894          4.011432        10.722232
1            -1.822208          4.052539        10.372803
2            -2.090110          4.147257        10.418942
3            -2.068434          4.045043        10.656953
4            -2.532326          3.969987        10.409853
..
783           ...           ...
784           ...           ...
785           ...           ...
786           ...           ...
787           ...           ...

    log_median_household_income  log_asthma  pred_asthma
0                  11.148506    2.424803    10.519589
1                  10.812532    2.415914    11.441303
2                  10.917250    2.322388    11.061450
3                  11.095454    2.322388    10.686823
4                  10.799922    2.360854    11.427801
..
783           ...           ...
784           ...           ...
785           ...           ...
786           ...           ...
787           ...           ...

```

[777 rows x 15 columns]

```

[46]: # Calculate error
model_gdf['error_asthma'] = model_gdf['pred_asthma'] - model_gdf['asthma']

# Set limit
abs_max = model_gdf['error_asthma'].abs().max()
clim = (-abs_max, abs_max)

# Plot error
model_plot = (
    gv.tile_sources.EsriImagery
    *
    gv.Polygons(

```

```

    model_gdf.to_crs(ccrs.Mercator()),
    vdims=['error_asthma', 'pred_asthma', 'asthma', 'tractID'],
    crs = ccrs.Mercator()
).opts(color = 'error_asthma', colorbar = True, tools = ['hover'],
      clim = clim, cmap = 'RdBu' # Diverging color map)
).opts(width = 600, height = 600, xaxis = None, yaxis = None, title = 'Asthma ↴
↪Prediction Error for Chicago, IL Census Tracts',
      clabel = 'Lower values indicate under-prediction, higher values indicate ↵
↪over-prediction')
)
model_plot

```

[46]: :Overlay  
.WMTS.I :WMTS [Longitude, Latitude]  
.Polygons.I :Polygons [Longitude, Latitude]  
(error\_asthma,pred\_asthma,asthma,tractID)

## 1.4 R2 Calculation

[34]: y\_true = model\_gdf['asthma'].values # observed  
y\_pred = model\_gdf['pred\_asthma'].values # predicted

[36]: # Example if you know which rows were training/testing  
train\_mask = model\_gdf.index.isin(X\_train.index)  
test\_mask = model\_gdf.index.isin(X\_test.index)

r2\_train = r2\_score(y\_true[train\_mask], y\_pred[train\_mask])  
r2\_test = r2\_score(y\_true[test\_mask], y\_pred[test\_mask])

print(f"Training R<sup>2</sup>: {r2\_train:.3f}")  
print(f"Test R<sup>2</sup>: {r2\_test:.3f}")

Training R<sup>2</sup>: 0.522

Test R<sup>2</sup>: 0.584

[39]: # Calculate R<sup>2</sup> for all data  
r2 = r2\_score(y\_true, y\_pred)  
print(f"R<sup>2</sup> (all): {r2:.3f}")

R<sup>2</sup> (all): 0.546

[38]: # Calculate R<sup>2</sup> for training and test sets  
train\_mask = model\_gdf.index.isin(X\_train.index)  
test\_mask = model\_gdf.index.isin(X\_test.index)

r2\_train = r2\_score(y\_true[train\_mask], y\_pred[train\_mask])  
r2\_test = r2\_score(y\_true[test\_mask], y\_pred[test\_mask])

```
print(f"Training R2: {r2_train:.3f}")
print(f"Test R2: {r2_test:.3f}")
```

Training R<sup>2</sup>: 0.522  
Test R<sup>2</sup>: 0.584

```
[48]: descriptor_text = (
    f"Model performance summary:\n"
    f"Using log_fract_veg, log_mean_patch_size, and log_mean_household_income,\n"
    f"\n"
    f"this model estimates log_asthma with R2 = {r2:.3f}."
)

text_panel = hv.Text(
    0.5, 0.5, descriptor_text
).opts(
    width=600,
    height=120,
    xaxis=None,
    yaxis=None,
    text_align='center',
    text_baseline='middle',
    text_font_size='10pt'
)
```

```
[49]: final_plot = (model_plot + text_panel).cols(1)
final_plot
```

```
[49]: :Layout
    .Overlay.I :Overlay
        .WMTS.I      :WMTS      [Longitude, Latitude]
        .Polygons.I :Polygons   [Longitude, Latitude]
(error_asthma,pred_asthma,asthma,tractID)
    .Text.I      :Text      [x,y]
```

```
[50]: # save final plot
hv.save(final_plot, 'descrip_asthma_error_map.html', fmt='html')
```